System Info

This document covers cross-cutting behavior, shared components, and infrastructure that span multiple pages of the resmon app. Page-specific behavior lives in the ten companion documents (`dashboard_info.md` , `deep_dive_info.md` , `deep_sweep_info.md` , `routines_info.md` , `calendar_info.md` , `results_and_logs_info.md` , `configs_info.md` , `monitor_info.md` , `repos_and_api_keys_info.md` , `settings_info.md`) and is referenced here rather than repeated.

System-Wide Overview

App Purpose

resmon (Research Monitor) is a local-first literature-surveillance desktop application that queries a curated catalog of open-access scholarly repositories, normalizes and deduplicates results across sources, persists them to a local SQLite database, optionally summarizes each paper via a user-configured BYOK LLM provider, and surfaces the resulting artifacts through an Electron + React UI. It supports on-demand manual runs ("Deep Dive" single-repo, "Deep Sweep" multi-repo) as well as cron-scheduled automated sweeps ("Routines") driven by a persistent APScheduler job store. Optional integrations include transactional email (SMTP), Google Drive artifact backup, and a closed-beta multi-user cloud mirror.

High-Level Architecture

The application is a three-process composition:

1. **Electron main process** — launches the Python backend as a child process, hosts the renderer window, bridges OS-level capabilities (directory chooser, shell open/reveal, cloud-auth keychain bridge via `window.resmonAPI` , see `frontend/src/api/client.ts`), and can attach to an already-running headless daemon instead of spawning its own backend (see `implementation_scripts/daemon.py`).
2. **Python FastAPI backend** (`resmon_scripts/resmon.py`) — exposes a local REST API on `127.0.0.1:<port>` (default `8742` , surfaced to the renderer via `window.resmonAPI.getBackendPort()`), owns the SQLite connection, orchestrates the sweep pipeline, hosts the APScheduler instance, and streams progress events over Server-Sent Events.
3. **Optional cloud service** — separate microservice (IMPL-27..40) backing the closed-beta multi-device sync surface; the desktop app consumes it through the `cloudClient` wrapper plus `useCloudSync` hook.

Frontend ↔ backend communication is exclusively `fetch` -over-localhost JSON plus SSE for progress. There is no IPC bridge between the renderer and Python; the Electron preload only exposes filesystem, shell, and cloud-auth helpers via a contextBridge-exposed `window.resmonAPI` surface.

Phase and Version Status

- `APP_NAME = "resmon"` , `APP_VERSION = "1.0.0"` (`implementation_scripts/config.py`).
- Phase 2 implementation is complete through IMPL-1..40, IMPL-R1..R12 (routine wiring + concurrent-execution monitor), and IMPL-AI1..AI14 (AI summarization wiring + provider expansion). This info-documentation tranche (IMPL-INFO1..INFO11) closes Phase 2; Phase 3 (Conclusion) is entered only under explicit user authorization.
- Schema version is tracked by `SCHEMA_VERSION` / `get_schema_version` in `implementation_scripts/database.py` .

Operational Modes

The app supports three operational modes that are orthogonal to the page surface:

1. **Local-only** — default. All routines, executions, and credentials live in the user's local SQLite database and OS keyring; cloud auth is not required.
2. **Hybrid** — local execution plus opt-in Google Drive artifact backup (see `Settings → Cloud Storage` panel and `settings_info.md`) and/or opt-in multi-device cloud sync (ingests remote `cloud_routines` / `cloud_executions` mirror rows via `useCloudSync`).
3. **Cloud-executed routines** (IMPL-37) — individual routines may be flagged `execution_location = "cloud"` so they run against the cloud service's worker path; local mirrors are pulled back on sync.

The Electron shell additionally supports a **headless daemon split** (IMPL-25): the backend can run as a long-lived OS service (launchd on macOS, `systemd --user` on Linux, Task Scheduler on Windows; see `implementation_scripts/service_manager.py`), and multiple renderer windows can attach to that single daemon via the lock file written by `daemon.py` .

Artifact Topology

All user-facing outputs land under `resmon_reports/` at the project root (`REPORTS_DIR` in `config.py`):

- `resmon_reports/markdowns/` — Markdown report per execution (authored by `report_generator.save_report`).
- `resmon_reports/pdfs/` — optional PDF exports (when the user exports a report; see `results_and_logs_info.md`).
- `resmon_reports/figures/` — static figures included in reports.
- `resmon_reports/latex/` — LaTeX renditions of reports.
- `resmon_reports/logs/` — per-execution logs emitted by `TaskLogger` plus the rotating app log (`resmon.log` , 5 MB × 3 backups; `logger.setup_app_logger`).
- `resmon_reports/info_docs/` — the 11 implementation-grounded info documents produced by this tranche.
- `resmon_experiments/` and `resmon_printouts/` — sibling folders used for ad-hoc experimental artifacts outside the report pipeline.

The SQLite database file itself lives at `DEFAULT_DB_PATH = PROJECT_ROOT / "resmon.db"` , and the OS-scoped state directory (lock file, daemon logs, service-unit-managed logs) is resolved by `daemon.state_dir()` (`~/Library/Application Support/resmon` on macOS, `$XDG_STATE_HOME/resmon` on Linux, `%LOCALAPPDATA%\resmon` on Windows).

Frontend

Electron Shell and React Runtime

The desktop surface is an Electron shell hosting a React 18 + TypeScript single-page application bootstrapped from `frontend/src/index.tsx` and rooted at `App.tsx` . The Electron main process exposes a narrow API surface to the renderer through a contextBridge-injected `window.resmonAPI` object (typed in `api/client.ts`) whose members include `getBackendPort` , `platform` , `versions` , `chooseDirectory` , `openPath` , `revealPath` , and the `cloudAuth` bridge used by `AuthContext` . The renderer never invokes Python directly — it always goes through the REST client.

HashRouter and Route Map

Routing is centralized in `App.tsx` via `react-router-dom` `HashRouter` . The route map is:

Path	Component	Info doc
`/`	`DashboardPage`	dashboard_info.md
`/dive`	`DeepDivePage`	deep_dive_info.md
`/sweep`	`DeepSweepPage`	deep_sweep_info.md
`/routines`	`RoutinesPage`	routines_info.md
`/calendar`	`CalendarPage`	calendar_info.md
`/results`	`ResultsPage`	results_and_logs_info.md
`/configurations`	`ConfigurationsPage`	configs_info.md
`/monitor`	`MonitorPage`	monitor_info.md
`/repositories`	`RepositoriesPage`	repos_and_api_keys_info.md
`/settings/*`	`SettingsPage` (nested router)	settings_info.md

`HashRouter` is used rather than `BrowserRouter` because the renderer is loaded from `file://` in the packaged Electron app, where push-state history is unreliable. The `Sidebar` , `Header` , and `MainContent` layout components wrap every route; the `FloatingWidget` monitor overlay sits outside `MainContent` so it survives route transitions.

Shared Contexts

Two React contexts wrap the entire route tree (see `App.tsx`):

- **`AuthProvider`** (`context/AuthContext.tsx`) — exposes `{ isSignedIn, email, accessToken, signIn, signOut, refreshAccessToken }` for the resmon-cloud account (IMPL-30). Access tokens are held only in the module-scoped `authStore` — never in React state, `localStorage` , or disk.

Refresh tokens live in the OS keychain, managed by Electron-main via the Python `keyring` bridge (service `resmon`, account `cloud_refresh_token`). `AuthContext` also tracks `sync_state` reported by the Electron bridge.

- **`ExecutionProvider`** (`context/ExecutionContext.tsx`) — global store for every concurrently-running execution. Keyed by execution id, it tracks `{ executionId, executionType: 'deep_dive' | 'deep_sweep' | 'automated_sweep', repositories, startTime, events, status, currentRepo`. It owns the SSE subscriptions, drives the FloatingWidget pulse, raises browser Notifications based on `notify_manual` / `notify_automatic_mode` (see `settings_info.md` → Notifications Panel), and exposes a `verboseLogging` toggle persisted under `localStorage["resmon.verboseLogging"]`. The multi-execution rewrite (IMPL-R8) replaced the single-execution store with `Record<number, ActiveExecution>` plus a focused-execution pointer that the FloatingWidget and `MonitorPage` tab strip read from.

Shared Hooks

Three cross-cutting hooks live in `frontend/src/hooks/`:

- **`useExecutionsMerged`** — fetches `/api/executions/merged`, which returns local SQLite executions plus cloud-mirror rows sorted by `start_time` DESC. Each row carries `execution_location ∈ {"local", "cloud"}` so `ResultsPage` can render the Local/Cloud badge and filter chip without inspecting id shapes. Auto-refreshes on: explicit `refresh()`, any `completionCounter` change from `ExecutionContext`, and the `resmon:cloud-sync-applied` window event emitted by `useCloudSync`.
- **`useRepoCatalog`** — loads `{catalog, bySlug, presence}` in parallel from the repository-catalog and credential-presence endpoints and exposes `refreshPresence()`. Used by the Deep Dive, Deep Sweep, Routines, and Repositories pages for a single source of truth on repository metadata and credential presence.
- **`useCloudSync`** — when the user is signed in and the window is focused, polls `GET /api/v2/sync?since=<last_synced_version>` on a 60-second interval, POSTs returned pages to the local daemon's `/api/cloud-sync/ingest`, drains `has_more` chains within a tick (cap `MAX_PAGES_PER_TICK = 50`), and triggers an immediate sync on focus-gain. On sign-out, it calls `/api/cloud-sync/clear` once. All errors are surfaced via `state.lastError` rather than propagated, per the "transient cloud outage must not crash the renderer" invariant.

API-Client Wrappers

Four API-client modules centralize all HTTP access:

- **`api/client.ts`** — default wrapper for the local FastAPI daemon. Resolves the base URL to `http://127.0.0.1:<port>` where the port comes from the Electron-injected `window.resmonAPI.getBackendPort()` (default `8742`). Every request sets `Cache-Control: no-store` and `Pragma: no-cache`. On non-2xx responses, the wrapper unwraps FastAPI's `{detail: "..."}` envelope and re-throws with a human-readable `"${status} ${statusText}: ${detail}"` message.
- **`api/cloudClient.ts`** — wrapper for the remote cloud service. Exports `CloudUnauthorizedError` so `useCloudSync` can distinguish token-expiry from transport errors.
- **`api/repositories.ts`** — typed helpers over the repository-catalog and credential-presence endpoints, consumed by `useRepoCatalog`.
- **`api/authStore.ts`** — module-scoped access-token store with refresh registration and a sign-out broadcaster.

FloatingWidget and Monitor Integration

The `FloatingWidget` (`components/Monitor/FloatingWidget.tsx`) is mounted once at the App root and renders an overlay that summarizes the focused execution plus a stack popover for any additional concurrent executions (IMPL-R10). It is driven entirely by `ExecutionContext` and reacts to `location.pathname` changes: it auto-minimizes on `/monitor`, any `/results`, and `/calendar`, since those pages already surface the full execution state on-screen. It supports cancel, clear, and focus-switch actions, and pulses on completion until the user acknowledges. The dedicated `/monitor` page renders the same data in a tab-strip + detail-pane layout; see monitor_info.md.

Theming and Badge Palette Conventions

Global stylesheet lives in `frontend/src/styles/` and is imported from `index.tsx`. Conventional class families used across pages:

- Status pill classes: `status-running`, `status-completed`, `status-failed`, `status-cancelled`, `status-cancelling`, `status-pending`.
- Repository-status glyphs in the widget/monitor: ✓ (done), ⟳ (querying), ✗ (error), ○ (pending) (`FloatingWidget.repoIcon`).
- Execution-type labels: `"Deep Dive"`, `"Deep Sweep"`, `"Auto Sweep"` (`FloatingWidget.typeLabel`).
- Stage labels used in the widget + Monitor page: `querying` → "Querying repositories", `dedup` → "Deduplicating results", `linking` → "Linking documents", `reporting` → "Generating report", `summarizing` → "AI summarization", `finalizing` → "Finalizing".
- The Results page adds a Local/Cloud origin badge driven by the `execution_location` field on merged rows.

Per-panel navigation links use the `tab-btn` / `tab-active` pair; see `settings_info.md` for the Settings tab-strip contract.

Cross-Page References

Cross-page flows that do not belong to a single info doc:

- Clicking a **Configuration** row on `/configurations` pre-fills the query form on `/dive` or `/sweep` (see `configs_info.md`).
- Completing an execution on `/dive`, `/sweep`, or the Routines "Run Now" action auto-navigates (or offers navigation) to `/results?execution_id=…` once the SSE stream closes (see `deep_dive_info.md`, `deep_sweep_info.md`, `routines_info.md`).
- The `/calendar` page derives its events from the same `/api/executions/merged` stream that feeds `/results`, keyed by `start_time` (`calendar_info.md`).
- The **credential presence map** exposed by `/api/credentials` is consumed by the Settings → Email/AI panels, the Repositories page, and the inline "Use a key just for this run" affordance on the Dive/Sweep/Routines forms.

Backend

FastAPI Application Shell

`resmon.py` constructs a single FastAPI application at module load: `app = FastAPI(title=APP_NAME, version=APP_VERSION)`. A single shared `sqlite3.Connection` (`_shared_conn`) is reused for every request via `_get_db()` / `_close_db()`; the DB path is overridable for tests via `_db_path`. `init_db` runs lazily on first access through `_db_initialized`. The database file defaults to `DEFAULT_DB_PATH` but tests can set `_db_path = ":memory:"` (with the scheduler jobstore redirected to a temp file so APScheduler's SingletonThreadPool can share schema).

Cross-Cutting Middleware

Two ASGI middlewares are registered, in order:

1. **`PrivateNetworkMiddleware`** — a raw ASGI middleware (not `BaseHTTPMiddleware`, because the latter buffers streaming responses). It injects the `Access-Control-Allow-Private-Network: true` header on every HTTP response start so Chromium's Private Network Access policy allows the `file://` renderer to reach `127.0.0.1`. SSE responses are unaffected because no buffering is introduced.
2. **`CORSMiddleware`** — `allow_origins=["*"]`, `allow_methods=["*"]`, `allow_headers=["*"]`. Safe because the backend binds to the loopback interface.

Admission Controller

The `admission` singleton (`implementation_scripts/admission.py`, `ExecutionAdmissionController`, ADQ-R3, IMPL-R1) is the single global gate on concurrent executions. Key invariants:

- Clamps `max_concurrent` to `[1, 8]` and `queue_limit` to `[1, 64]`; defaults `DEFAULT_MAX_CONCURRENT = 3`, `DEFAULT_QUEUE_LIMIT = 16`.
- **Manual admission** is reject-or-pass: if no slot is free, the REST layer raises HTTP 429 (see Deep Dive / Deep Sweep / manual-search endpoints via IMPL-R2).
- **Routine admission** falls through to a bounded FIFO queue (`deque`) that drains as slots free. Overflow past `queue_limit` is dropped with a logged error (APScheduler worker threads must never raise).
- All state transitions are guarded by `_lock`. Queue drains spawn a fresh daemon thread per dispatch so the `finally`-path of a completing pipeline is never blocked by the next pipeline starting.
- Live values are reloaded from the `settings` table on startup (`_init_admission_on_startup` → `_hydrate_admission_from_db`) and updated via `PUT /api/settings/execution` (see `settings_info.md` → Advanced Panel and IMPL-R12).

SweepEngine Pipeline

`SweepEngine` (`implementation_scripts/sweep_engine.py`) orchestrates the complete query → normalize → dedup → report pipeline for both manual and routine-fired executions. The same code path is entered via `execute_dive` (single repo), `execute_sweep` (multi-repo), or `prepare_execution` + `run_prepared` (split pattern used when the REST layer needs the exec id before kicking off the background thread). Per-execution stages emit `progress_store` events at `querying`, `dedup`, `linking`, `reporting`, `summarizing`, and `finalizing` — the same identifiers consumed by the FloatingWidget's `stageLabel` map.

Cross-cutting responsibilities:

- **Required-credentials gate** — the `_REQUIRED_CREDENTIALS` table maps repositories that cannot return any results without a key (`core`, `ieee`, `nasa_ads`, `springer`) to their credential name. Missing credentials short-circuit the repo with a user-visible error.

- **Cooperative cancellation** — `_search_with_heartbeat` polls `progress_store` 's cancel flag on a 2-second heartbeat; `_ExecutionCancelled` is caught by `_handle_cancellation` to persist a partial report.
- **Optional auto-backup** — `_maybe_auto_backup` invokes `cloud_storage.upload_directory` when the `cloud_auto_backup` setting is true and a Google Drive token is present.
- **Per-execution task log** — `TaskLogger` writes structured per-execution logs under `resmon_reports/logs/` (see `results_and_logs_info.md`).

ResmonScheduler

`ResmonScheduler` (`implementation_scripts/scheduler.py`) wraps APScheduler's `BackgroundScheduler` with a `SQLAlchemyJobStore` backed by `sqlite:///<DEFAULT_DB_PATH>` (ADQ-3). Jobs persist across app restarts. The module is deliberately decoupled from FastAPI and `SweepEngine` through a dispatcher indirection (IMPL-R3):

- `set_dispatcher(fn)` installs the callback invoked on each routine fire. `fn` receives `(routine_id: int, parameters: str)` .
- When no dispatcher is installed, `_routine_callback` logs an error and returns — the APScheduler worker thread never raises.
- Cron parsing is strict 5-field (`minute hour day month day_of_week`) via `_parse_cron` .

The FastAPI startup hook (IMPL-R4) calls `set_dispatcher(_dispatch_routine_fire)` , instantiates a single `ResmonScheduler` , starts it, then walks `get_routines(conn)` and re-registers every active routine so jobs survive a cold restart. The shutdown hook gracefully stops the scheduler and clears the dispatcher. `_dispatch_routine_fire` consults the admission controller before spawning a pipeline, stamps `executions.routine_id` , launches via `_launch_execution` , and updates `routines.last_executed_at` on completion. See routines_info.md for the full fire lifecycle and settings_info.md → Advanced Panel for the `/api/scheduler/jobs` diagnostics surface.

ProgressStore and SSE Progress Stream

`ProgressStore` (`implementation_scripts/progress.py`) is the thread-safe in-memory bus between the pipeline worker threads and the SSE endpoint. Each `exec_id` owns its own event list, `threading.Lock` , `threading.Event` cancel flag, and completion boolean. The API is `register(exec_id)` , `emit(exec_id, event)` , `get_events(exec_id, since)` , `mark_complete(exec_id)` , `cleanup(exec_id)` . The SSE endpoint pages through `get_events(exec_id, since=<cursor>)` until `_completed` flips; post-completion the REST layer persists the buffered events into the `execution_progress` table via `save_progress_events` so `ReportViewer` can reconstruct the stream after the process restarts (IMPL-21). See monitor_info.md and results_and_logs_info.md.

Cloud Integration

Two distinct cloud surfaces exist and must not be conflated:

1. **Google Drive artifact backup** (`implementation_scripts/cloud_storage.py`) — least-privilege `drive.file` OAuth 2.0 scope, `InstalledAppFlow.from_client_secrets_file` using `credentials.json` at the project root, token serialized and stored in the OS keyring under `google_drive_token` . Public helpers: `check_connection` , `authorize_google_drive` , `revoke_authorization` , `upload_directory` , `is_token_stored` , `probe_api` . Surfaced through `Settings → Cloud Storage` (see settings_info.md).
2. **resmon-cloud account + multi-device sync** (IMPL-27..40) — separate microservice backing the closed-beta experience. Desktop consumes it via `AuthContext` , `api/cloudClient.ts` , `useCloudSync` , and the local-daemon ingest endpoint (`/api/cloud-sync/ingest` , `/api/cloud-sync/clear`). Envelope-encrypted per-user credentials (ADQ-11, IMPL-31): per-user DEK wrapped by a KMS-held KEK. JWKS-verified JWTs, per-user row-level security (IMPL-28..29), signed-URL artifact fetches (IMPL-34), per-user rate limiting (IMPL-33). The `Settings → Cloud Account` panel is currently "under construction" — backend endpoints exist but the UI does not wire up a hosted IdP in this build (see settings_info.md).

Credential Manager and Ephemeral Key Stack

`implementation_scripts/credential_manager.py` is the only module that touches the OS keyring (ADQ-4). It defines two frozensets used as whitelists by `PUT /api/credentials/{name}` :

- `AI_CREDENTIAL_NAMES` — `openai_api_key` , `anthropic_api_key` , `google_api_key` , `xai_api_key` , `meta_api_key` , `deepseek_api_key` , `alibaba_api_key` , `custom_llm_api_key` (ADQ-AI9).
- `SMTP_CREDENTIAL_NAMES` — `smtp_password` .

Additional repository credentials (`core_api_key` , `ieee_api_key` , `nasa_ads_api_key` , `springer_api_key` , plus any optional/recommended keys from the catalog) are whitelisted via `repo_catalog.credential_names` . Core API: `store_credential` , `get_credential` , `delete_credential` , `validate_api_key` . The ephemeral-key stack (`push_ephemeral` , `pop_ephemeral` , `get_credential_for`) lets Deep Dive / Deep Sweep / Routines pass a per-execution key that lives only in-process, is keyed by `exec_id` , is never logged, and takes precedence over the persisted keyring value (IMPL-23). Credential values are never returned by any GET endpoint; only presence booleans are exposed through `/api/credentials` .

LLM Factory and Provider Whitelist

`implementation_scripts/llm_factory.py` (`build_llm_client_from_settings`) turns the persisted `ai_*` settings into a concrete `RemoteLLMClient` or `LocalLLMClient` , or returns `None` when the provider is unset or its credentials are missing (ADQ-AI7 / F6 — "AI unconfigured" is a silent no-op, never an exception). The only `ValueError` the factory raises is for `ai_provider == "custom"` with an insecure `ai_custom_base_url` : `_validate_custom_base_url` rejects non-HTTPS URLs unless the host is a loopback (`localhost` , `127.0.0.1` , `::1`). API keys are never logged, never included in exception messages, and never returned by any public function (ADQ-AI8; OWASP A02).

Key lookup order per ADQ-AI9: ephemeral `{provider}_api_key` → ephemeral `custom_llm_api_key` (custom provider only) → keyring `{provider}_api_key` → keyring `custom_llm_api_key` (custom provider only). The Settings → AI panel drives the same whitelist (IMPL-AI9 / AI10), the custom-URL guard is enforced both in the UI (IMPL-AI12) and in this factory, and `implementation_scripts/ai_models.list_available_models` (surfaced through the AI panel's "Load models" button) uses `httpx` to fetch the provider's model index without logging credentials. See settings_info.md → AI Panel.

Database Schema and Cloud-Mirror Cache

The SQLite schema is owned by `implementation_scripts/database.py` (ADQ-5); all queries are parameterized. Core tables referenced across pages (not exhaustive):

- `routines` — user-defined automation; `execution_location ∈ {"local","cloud"}` (IMPL-37).
- `executions` — every manual or routine-fired run; `routine_id` stamped by the dispatcher.
- `execution_progress` — persisted progress events (IMPL-21).
- `documents` + cross-tables — deduplicated paper records and per-execution links.
- `configurations` — saved query presets (`configs_info.md`).
- `settings` — flat key/value settings table accessed via `get_setting` / `set_setting` .
- `cloud_routines` / `cloud_executions` — cursor-synced read-only mirror of the cloud service's rows (IMPL-35..36).
- `sync_state` — holds `last_synced_version` per `get_last_synced_version` / `set_last_synced_version` .
- Cloud-cache bookkeeping — `record_cloud_cache_entry` , `touch_cloud_cache_entry` , `get_cloud_cache_entry` , `get_cloud_cache_total_bytes` , `evict_cloud_cache_if_needed` , with `CLOUD_CACHE_MAX_BYTES_DEFAULT` as the LRU bound.

Schema-version tracking (`SCHEMA_VERSION` , `get_schema_version`) runs migrations on startup via `init_db` .

Settings Surface Summary

Every persisted preference is eventually a row in the `settings` table. The Settings page is divided into eight panels (see settings_info.md): **Email**, **Cloud Account**, **Cloud Storage**, **AI**, **Storage**, **Notifications**, **Advanced**, **About App**. Cross-cutting behaviors reached from multiple pages:

- `notify_manual` / `notify_automatic_mode` — consumed by `ExecutionContext.maybeNotifyCompletion` .
- `cloud_auto_backup` — consumed by `SweepEngine._maybe_auto_backup` .
- `max_concurrent_executions` / `routine_fire_queue_limit` (IMPL-R12) — consumed by the admission controller.
- `ai_provider` / `ai_model` / `ai_local_model` / `ai_custom_base_url` / `ai_custom_header_prefix` / `ai_summary_length` / `ai_tone` / `ai_temperature` / `ai_extraction_goals` — consumed by `llm_factory.build_llm_client_from_settings` and the summarizer.
- `export_directory` — consumed by report export actions on the Results page.
- `pdf_policy` / `txt_policy` / `archive_after_days` — reserved for a future per-paper artifact feature; currently no effect on the sweep pipeline (documented as such in `settings_info.md`).

Logging Topology

Three logging surfaces coexist:

1. **Rotating app log** — `logger.setup_app_logger(log_dir)` creates a `RotatingFileHandler` at `resmon_reports/logs/resmon.log` (5 MB × 3 backups, UTF-8). Format: `%(asctime)s [%(levelname)s] %(name)s: %(message)s` . Every module uses `logging.getLogger(__name__)` so log lines are traceable to their source.
2. **Per-execution task log** — `TaskLogger` writes a structured log per execution id into `resmon_reports/logs/` and is linked from `ReportViewer` (see results_and_logs_info.md).
3. **In-memory progress events** — `progress_store` (see above) — transient, then persisted to `execution_progress` on completion.

The renderer-side `verboseLogging` toggle (stored under `localStorage["resmon.verboseLogging"]`) controls only the client-side console verbosity of `ExecutionContext` ; it does not alter backend log levels.

Report Artifact Layout

Reports generated by `report_generator.generate_report` + `save_report` land under the directories enumerated in "Artifact Topology" above. File naming is deterministic per execution (stamp + execution id + execution type — see `resmon_implementation_guide.md` Appendix D for the canonical contract). Exports (PDF, LaTeX) are triggered from `ResultsPage` via explicit user actions; the Markdown report is always produced. See results_and_logs_info.md for the per-report viewer UI and for how `report_exporter.py` is invoked.

Cross-Page References

Cross-cutting backend flows that touch more than one page:

- **Notifications pipeline** — `GET/PUT /api/settings/notifications` is owned by the Settings → Notifications panel, but the actual notification firing happens inside `ExecutionContext.maybeNotifyCompletion` across every page that runs executions.
- **Execution-settings admission wiring** — `GET/PUT /api/settings/execution` lives in the Settings → Advanced panel (IMPL-R12) but binds the singleton used by Deep Dive, Deep Sweep, and routine dispatch.
- **Credential presence map** — `GET /api/credentials` returns a single `{name: bool}` envelope consumed simultaneously by `useRepoCatalog`, `EmailSettings`, `AISettings`, and the Repositories page's scope selector (IMPL-38, keyring-vs-cloud).
- **Merged executions feed** — `/api/executions/merged` is the sole backend source for both the Results page list and the Calendar page event stream.
- **Startup reconciliation** — `_init_admission_on_startup` and `_init_scheduler_on_startup` together guarantee that the admission cap and active-routine registrations match the persisted DB state before any request is served.

System Info

System-Wide Overview

App Purpose

High-Level Architecture

Phase and Version Status

Operational Modes

Artifact Topology

Frontend

Electron Shell and React Runtime

HashRouter and Route Map

Shared Contexts

Shared Hooks

API-Client Wrappers

FloatingWidget and Monitor Integration

Theming and Badge Palette Conventions

Cross-Page References

Backend

FastAPI Application Shell

Cross-Cutting Middleware

Admission Controller

SweepEngine Pipeline

ResmonScheduler

ProgressStore and SSE Progress Stream

Cloud Integration

Credential Manager and Ephemeral Key Stack

LLM Factory and Provider Whitelist

Database Schema and Cloud-Mirror Cache

Settings Surface Summary

Logging Topology

Report Artifact Layout

Cross-Page References